

Module 1

Microcontroller Programming Techniques

Textbook:

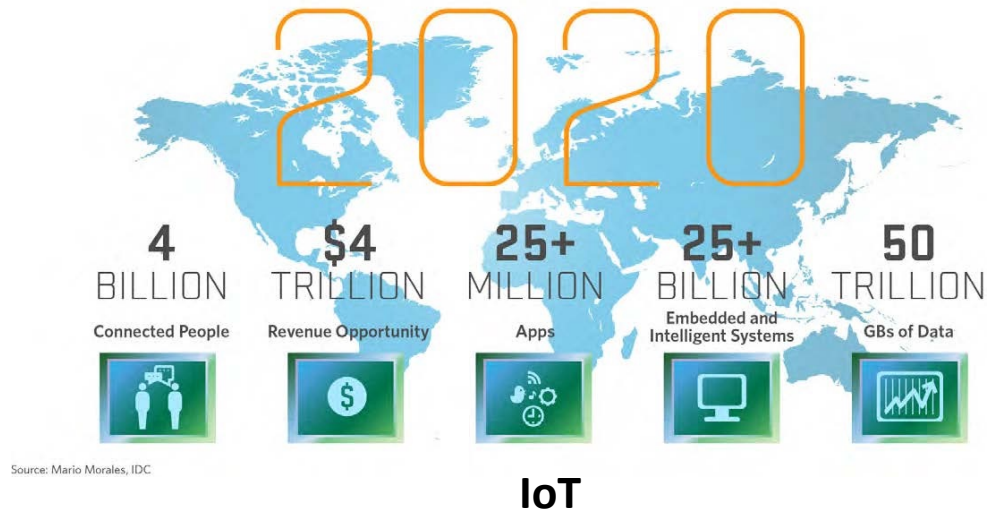
Chapter 1 (p1-26)

Chapter 2 (p27-54): Mostly review from 270

Learning Outcome #1

“An ability to program a microcontroller to perform various tasks”

Why?



Embedded systems in cars

Learning Outcome #1

“An ability to program a microcontroller to perform various tasks”

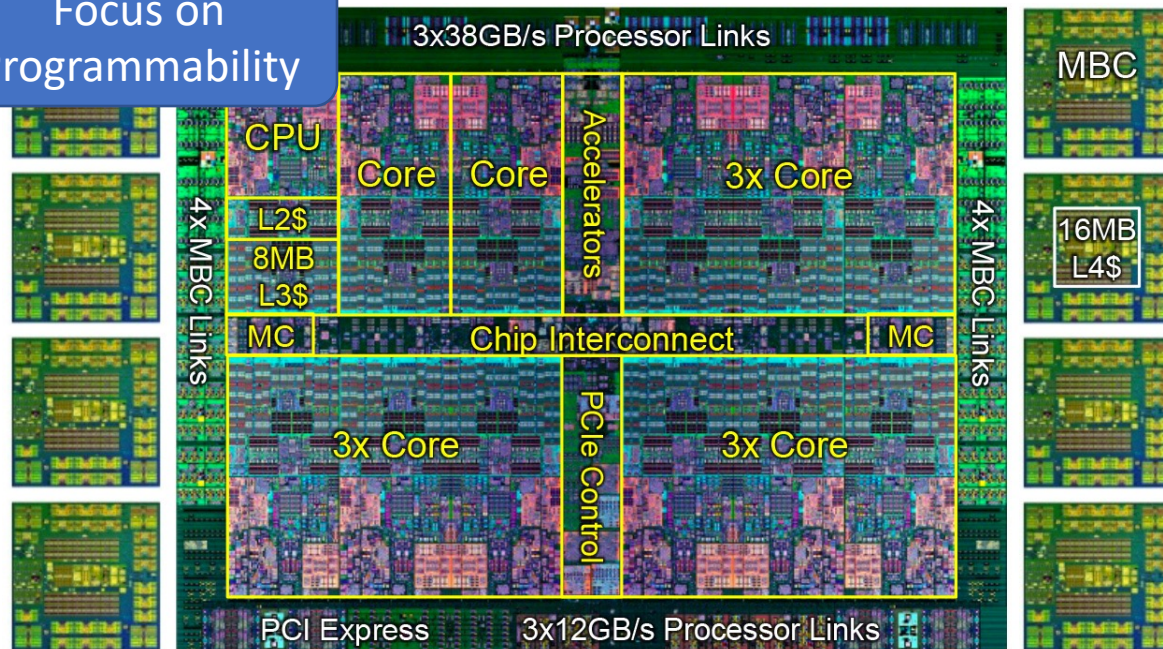
How?

- A. Architecture and Programming Model
- B. Instruction Set Basics
- C. Addressing modes and conditionals
- D. Practical Assembly Language Programming
- E. Stacks and Recursive Functions*

What makes a Microcontroller?

General Purpose

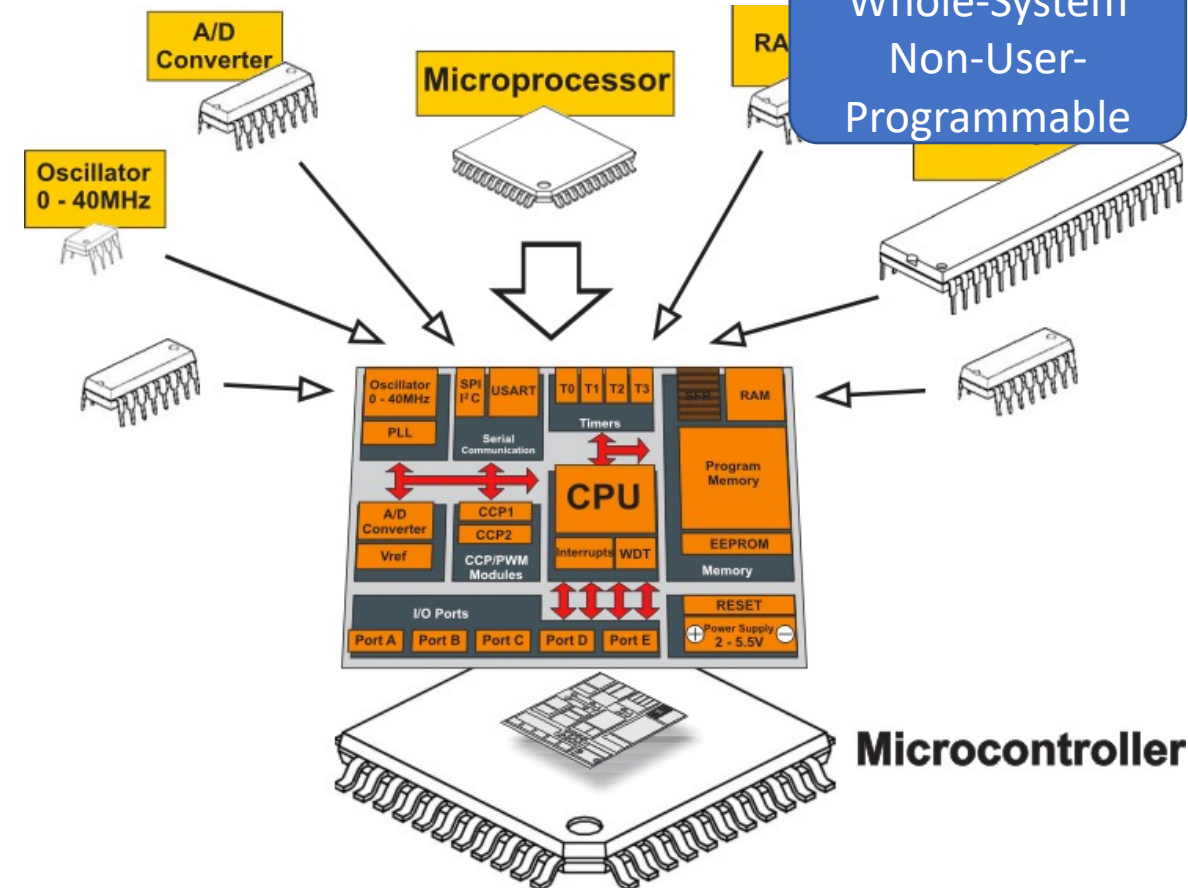
Focus on
Programmability



IBM Power 8 Microprocessor

Embedded

Whole-System
Non-User-
Programmable



Different Objectives

General Purpose

- High-performance
- Need to run “anything” relatively well
- Less energy constrained
- “Limitless” Memory
- Silicon devoted to compute and caches

Embedded

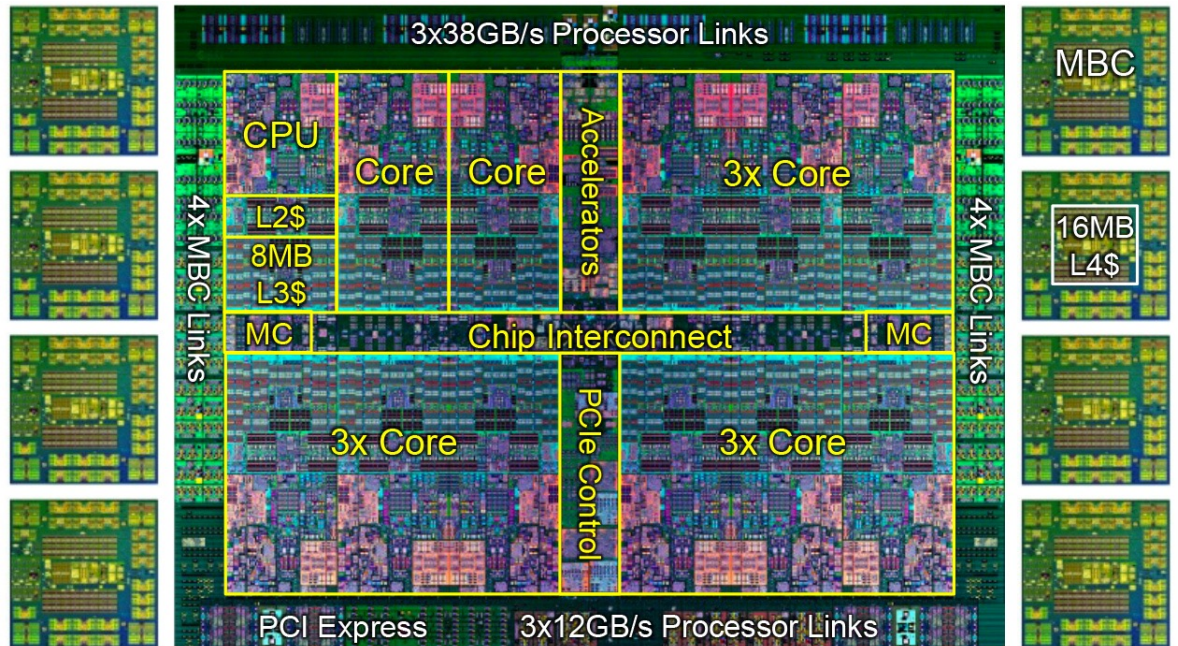
- Limited ability
- Inexpensive
- Limited memory
- Low-energy
- Silicon devoted to many peripherals (although core still consumes a lot)

Characteristics of general-purpose processors

- Virtual memory (Page Tables, TLBs)
- Deep cache hierarchy
- Large number of registers
- Hardware floating point
- Deep Pipelines/ Out of Order Exec.
- Complex Prediction Mechanisms
- Multicore


Take ECE 437!

Interrupts mostly get in the way




Characteristics of embedded processors

- Live and die by interrupts!
- Few registers (need to switch contexts fast)
- Circuitry devoted to analog
- Often rely on hand-tuned assembly code



This basics of all computers

- What does it do?
 - Fetch instruction from address in Program Counter (PC)
 - Execute instruction
 - Update PC to point to next instruction



The Basics

- C statement
- $f = (g + h) - (i + j)$

What does the CPU
need to do?

What assembly instructions might look like

- Assembly instructions*

add t0, g, h // $t0 = g + h$

add t1, i, j // $t1 = i + j$


sub f, t0, t1 // $f = t0 - t1$

- Opcodes/mnemonic, operands, source/destination

Not the only of doing
business!

Instruction Set Architectures (ISAs)

- Set of operations **visible to the programmer**
- ISA is the contract between the SW and HW




- Examples: CPU12, x86, PowerPC, **ARM**, SPARC, etc...

Different Types of ISAs


- CISC (Complex Instruction Set Computer)
- RISC (Reduced Instruction Set Computer)
- DSPs (Digital Signal Processors)

ARM Cortex-M0
(32-bits)

ISA defines Interaction between processor and memory




??



Some common Architectures

How to compute " $Z = X + Y$ ":

(b) Accumulator




Accumulator

```
LoadA  X
AddA   Y
StoreA Z
```

Reg/Reg


```
Load  R1,X
Load  R2,Y
Add   R3,R1,R2
Store R3,Z
```

(d) Register-register/load-store




ARM Cortex-M0
(ECE 362)

Memory Types



SRAM (Static Random Access Memory)

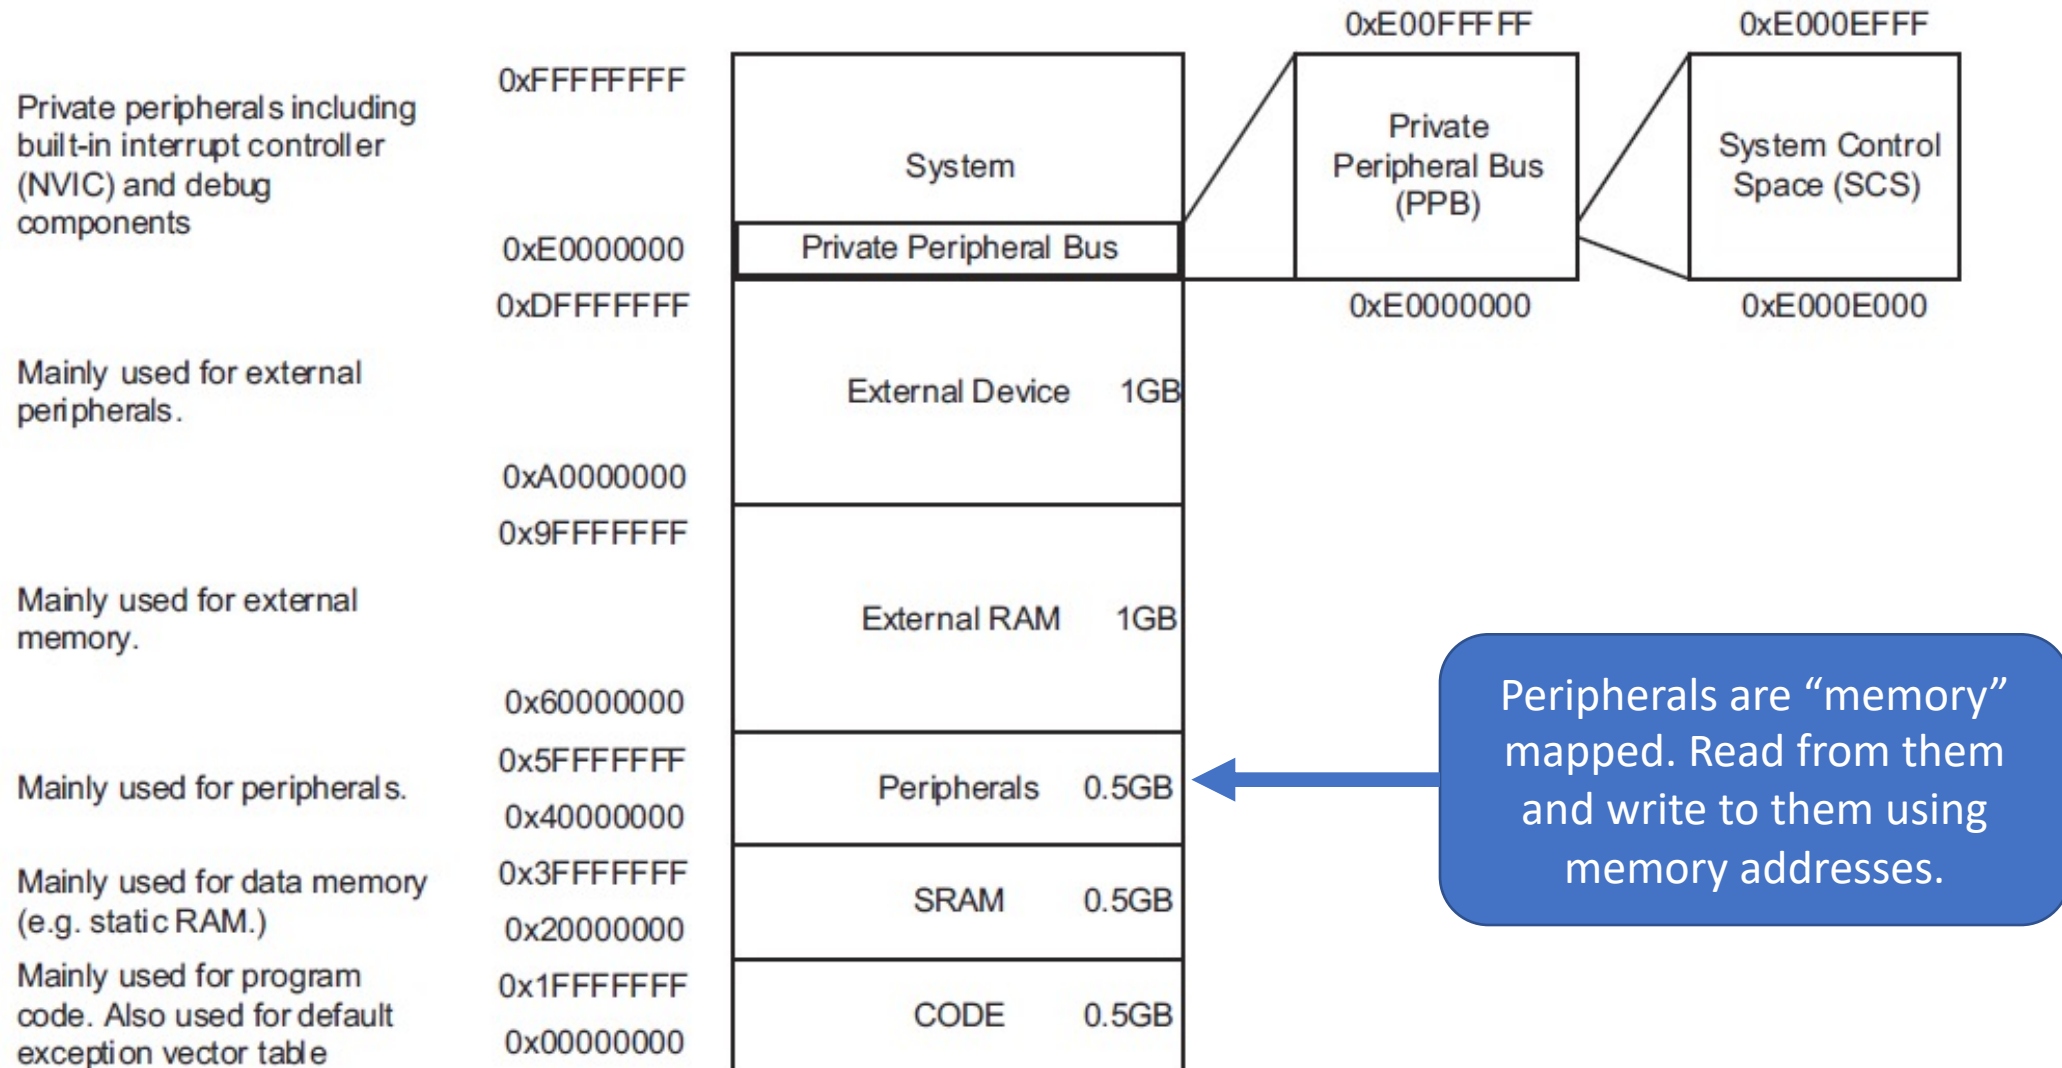
- Volatile
 - Variables
 - Stack
 - Buffers
 - Test Code



Flash Memory


- Non-Volatile
 - App Code
 - Static Data
 - Vectors (reset and interrupts)

Memory Map



Microcontroller Breakdown

- 32-bit ARM Cortex-M0 CPU.
- 8MHz clock multiplied x6 by a phase-locked loop (PLL) to get 48 MHz.
- 3-stage pipeline. Most instructions take 3 cycles, but a new instruction starts each cycle.
- Built-in Flash ROM to store program and SRAM to hold read/write data.
- 32 maskable interrupt channels w/ 4 priority levels
- 12-bit ADC, 12-bit DAC
- SPI, I2S, I2C, USART, CAN, HDMI CEC peripherals
- Several timers with integrated PWM
- So much more:
- <https://engineering.purdue.edu/ece362/refs/>



ARM Cortex-M0 CPU: Our Embedded CPU

- Maximum power dissipation: 400 – 500mW
- No heatsink needed. Easy to “embed.”
- Sometimes people try to embed a conventional high-performance CPU.
 - This kind of thing needs a heatsink, fan, ventilation, etc.
- By comparison, laptop CPUs use ~50W TDP, data center CPUs ~200W, Discrete GPUs >250W!
- Not good for embedded use.

Thumb ISA

- Used to improve code density
 - Cannot access r8 through r12
 - Implements fewer op-codes
- Still operate on 32-bit registers



Thumb Encoding



Encoding T1

All versions of the Thumb ISA.

ADDS <Rd>, <Rn>, #<imm3>

ADD<C> <Rd>, <Rn>, #<imm3>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	1	1	1	0	imm3			Rn			Rd		

Encoding T2

All versions of the Thumb ISA.

ADDS <Rdn>, #<imm8>

ADD<C> <Rdn>, #<imm8>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	1	1	0	Rdn			imm8							

Encoding T3

ARMv7-M

ADD{S}<C>.W <Rd>, <Rn>, #<const>


15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	1	1	1	0	i	0	1	0	0	0	S	Rn			0	imm3			Rd			imm8									

Encoding T4

ARMv7-M


ADDW<C> <Rd>, <Rn>, #<imm12>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	1	1	1	0	i	1	0	0	0	0	0	Rn			0	imm3			Rd			imm8									



Our ISA's Registers

- 16 registers
 - R15 is the PC.
 - R13 is the SP.
 - R14 is the "Link Register" (LR)
 - Most instructions can only use R0 – R7.
- Program Status Register (PSR).
 - Upper four bits are the flags.

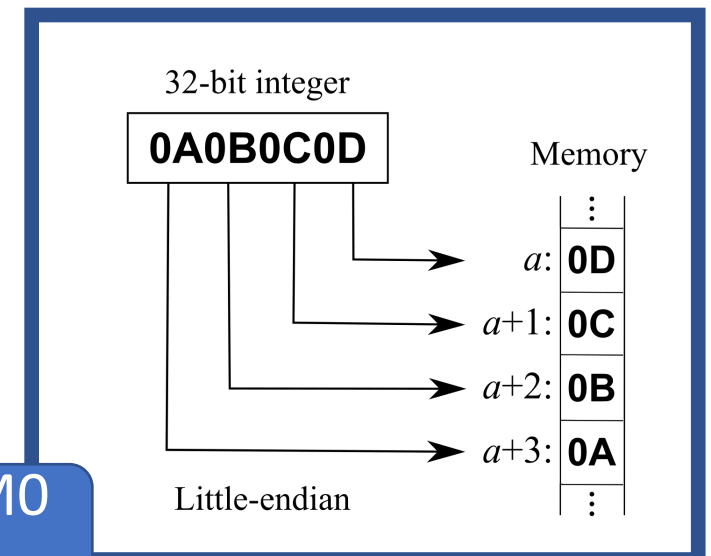



Upper PSR Bits

- “N” – “negative” flag (most significant bit (sign) of computation)
- “Z” – “zero” flag (set if result of computation is zero)
- “C” – “carry” flag. (set to 1 if the instruction results in a carry condition, for example an unsigned overflow on an addition)
- “V” – “overflow” flag (set to 1 if the instruction results in an overflow condition, for example a signed overflow on an addition)

Endianness

- Word storage
- Word = bytes (0x400,0x401,0x402,0x403)?
- Word = bytes (0x403,0x402,0x401,0x400)?
 - It depends...
- **Big endian**: MS Byte at address xxxxxx**00b**
 - e.g., IBM, SPARC
- **Little endian**: MS Byte at address xxxxxx**11b**
 - e.g., Intel x86
- Mode selectable
 - e.g., PowerPC, MIPS, ARM (generally)



ARM Cortex-M0
(ECE 362)

Instruction Representation

- Just a bunch of bits
- Remember? Simple Computer?

Opcode	Mnemonic	Function Performed
0 0 0	LDA <i>addr</i>	Load A with contents of location <i>addr</i>
0 0 1	STA <i>addr</i>	Store contents of A at location <i>addr</i>
0 1 0	ADD <i>addr</i>	Add contents of <i>addr</i> to contents of A

Example Encoding for Our Instructions

A6.7.2 ADD (immediate)

This instruction adds an immediate value to a register value, and writes the result to the destination register. It updates the condition flags based on the result.

Encoding T1 All versions of the Thumb instruction set.

ADDS <Rd>, <Rn>, #<imm3>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	1	1	1	0	imm3			Rn			Rd		

d = UInt(Rd); n = UInt(Rn); setflags = !InITBlock(); imm32 = ZeroExtend(imm3, 32);

Encoding T2 All versions of the Thumb instruction set.

ADDS <Rdn>, #<imm8>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	1	1	0	Rdn			imm8							

d = UInt(Rdn); n = UInt(Rdn); setflags = !InITBlock(); imm32 = ZeroExtend(imm8, 32);

General Encoding format

A5.2 16-bit Thumb instruction encoding

The encoding of 16-bit Thumb instructions is:

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
opcode															

Table A5-1 shows the allocation of 16-bit instruction encodings.

Table A5-1 16-bit Thumb instruction encoding

opcode	Instruction or instruction class
00xxxx	<i>Shift (immediate), add, subtract, move, and compare</i> on page A5-85
010000	<i>Data processing</i> on page A5-86
010001	<i>Special data instructions and branch and exchange</i> on page A5-87
01001x	Load from Literal Pool, see <i>LDR (literal)</i> on page A6-141
0101xx	<i>Load/store single data item</i> on page A5-88
011xxx	
100xxx	
10100x	Generate PC-relative address, see <i>ADR</i> on page A6-115
10101x	Generate SP-relative address, see <i>ADD (SP plus immediate)</i> on page A6-111
1011xx	<i>Miscellaneous 16-bit instructions</i> on page A5-89
11000x	Store multiple registers, see <i>STM, STMIA, STMEA</i> on page A6-175
11001x	Load multiple registers, see <i>LDM, LDMIA, LDMFD</i> on page A6-137
1101xx	<i>Conditional branch, and Supervisor Call</i> on page A5-90
11100x	Unconditional Branch, see <i>B</i> on page A6-119

Addressing Modes

- Definition: The CPU uses an addressing mode to determine how instructions get operands.
- When getting those operands from memory, it specifies how the address is computed.

C-code

$f = (g + h) - (i + j)$

“assembly”

add $t0$, g , h // $t0 = g + h$

add $t1$, i , j // $t1 = i + j$

sub f , $t0$, $t1$ // $f = t0 - t1$

g, h, i, j and f all need to be a 32-bit address.
Addressing modes tell us how we compute that address.

Addressing Modes

- Purpose:
 - Need to access memory
 - Need an **ADDRESS**

??

Memory	
Addr	Value
0	255
1	0
2	270
65535	69

Value not possible

Simple Computer Addressing Modes

The Simple Computer had 3 addressing modes:

- **Register**: Operand is read from a register
 - ADD R6,R11
- **Immediate**: Operand is encoded directly in the instruction:
 - ADDI R5,#0f
- **Absolute**: Operand referred to a specific memory location.
 - LDL R3,01fc

Cortex-M0 ISA

- RISC with compromises
 - Most instructions can refer only to registers R0 – R7.
 - Only small constants are used in instructions.
 - Load/Store architecture: Memory references separate from arithmetic.
- Four major addressing modes:
 - **Immediate**: small value contained in instruction
 - **Register**: src/dst register in instruction
 - **Offset**: value is at an address that's a relative (constant) offset from a register
 - **Indexed**: value is at address specified by register plus a second register

Accessing Memory with Loads and Stores

- The Cortex-M0 CPU cannot reference memory with arithmetic or logical instructions.
- It can only reference memory with load and store instructions.
 - That's why it's called a load/store architecture.
- It also cannot refer to an absolute address.
 - Instead, it uses an address stored in a register.
 - A load or store is relative to one of registers R0 – R7.

The LDR (Load Register) Instruction

- Let's say R0 holds the value 0x20000000.
- We want to load the four-byte value contained in addresses
- 0x20000000 – 0x20000003 into register R1.
- This is the instruction:
 - LDR R1, [R0]
- “Use the value in R0 as an address. Read 4-bytes starting from that address, and put the result in R1.”
- **Caveat: R0 must be evenly divisible by 4.**
 - If not, it will invoke a fault handler. You don't want that.

The STR (Store Register) Instruction

- Let's say R0 holds the value 0x20000000.
- We want to store the four-byte value contained in R1 to the addresses 0x20000000 – 0x20000003.
- This is the instruction:
 - STR R1, [R0]
- “Use the value in R0 as an address. Write the value of R1 into the 4-bytes starting at that address.”
- **Caveat: R0 must be evenly divisible by 4.**
 - If not, it will invoke a fault handler. You don't want that.

C-code
f = (g + h)

“assembly”
add f, g, h

Cortex M-0
Assume R0 holds 0x0 (g)
R1 holds 0x4 (h)
R2 holds 0x8 (f)

LDR R3, [R0]
LDR R4, [R1]
ADDS R5, R3, R4
STR R5, [R2]

Memory

Addr Value

0	1
1	0
2	0
3	0
4	5
5	0
6	0
7	0